

The Braid-Accumulator Ledger: A Novel Consensus and Data Structure Model for Decentralized Systems

Author: K. Barnhart

Abstract

This paper introduces the Braid-Accumulator Ledger (BAL), a consensus and state model that departs fundamentally from both blockchains and directed acyclic graph (DAG) systems. Rather than establishing a global or partial order over individual transactions, BAL treats state changes as commutative, causally-scoped deltas inside independent fibers and commits them to cryptographic accumulators. Consensus executes only on periodic vector commitments of per-fiber accumulator roots via a Byzantine Fault Tolerant checkpoint. The result is a ledger where the principal object of agreement is membership in committed sets, not the relative sequence of events. We present the core architecture, analyze security and data-availability assumptions, outline compatibility with order-dependent subsystems, and compare BAL against blockchains, DAG protocols, and stateless-accumulator research. We argue that BAL enables horizontal scalability, simplified light-client verification, and reorg-resistant finality while exposing a distinct trade-off surface for developers and protocol designers.

1. Introduction

1.1 Motivation

Modern decentralized ledgers largely inherit the design choice that disputes and conflicts should be resolved by ordering: first come, first served. This choice has virtues - conceptual simplicity and intuitive replay semantics - but it couples all applications to a single sequencing resource. Even independent actions compete for order, creating congestion externalities, opening room for ordering-based value extraction, and forcing the system to pay a consensus cost that scales with transaction volume rather than the minimal information needed for agreement. BAL explores a different axis: design operations that commute by construction, and have consensus attest only to compact set commitments, not full sequences.

1.2 Problem Statement and Contributions

Blockchains impose a total order; DAGs impose a partial order; stateless designs compress proofs but preserve ordering; and CRDT-style ledgers commute operations but rarely bind them to open BFT consensus with a strong data-availability guarantee. We design a ledger that minimizes reliance on order while preserving safety, liveness, and verifiability at scale. Contributions: (i) a fibered state model with deterministic routing of deltas; (ii) dual

accumulators tracking both effect sets and spent-set uniqueness to preclude double claims; (iii) a periodic BFT checkpoint over a vector of fiber roots; (iv) a light-client proof system requiring only the latest checkpoint and local witnesses.

2. System Architecture

2.1 Fibers

Fibers are independently validated sub-ledgers mapped deterministically by a routing function $h(\text{asset}|\text{contract}|\text{key}) \rightarrow \text{fiber_id}$. They are not shards that require cross-shard consensus for correctness; rather they are accountability boundaries that maintain their own accumulators and invariants. The global system agrees only on their root commitments at checkpoints. This yields strong locality and horizontal concurrency.

2.2 Deltas

Deltas are signed operations that update a fiber's abstract state. They are designed to commute: examples include set add or remove with tombstones, counter increments or decrements within bounds, and escrow-style transfers with unique claim identifiers. Deltas carry witnesses establishing preconditions such as membership of an unspent capability or non-membership of a nonce, and postconditions such as inclusion of a new claim.

2.3 Accumulators and Spent-Sets

Each fiber maintains at least two accumulators: an effect accumulator for committed facts and a spent-set accumulator for uniqueness and replay protection. Candidate constructions include RSA accumulators, KZG polynomial commitments, and vector commitments. Witness updates can be batched; clients refresh proofs using DA-available deltas since their last sync.

2.4 Checkpoints and Proof Model

At fixed intervals T , validators finalize a vector $R = [r_1, r_2, \dots, r_F]$ of current fiber roots via BFT. Only the vector and metadata are finalized; raw deltas live in the data-availability layer. A client verifies a claim using the latest checkpoint, a membership or non-membership witness against the relevant fiber root, and a quorum signature.

3. Consensus Protocol

3.1 Network and Failure Model

We assume n validators with authenticated channels, partial synchrony, and at most $f < n/3$ Byzantine faults as in classical BFT. Validators maintain per-fiber mempools and tentative accumulator states. Messages are scoped by fiber to minimize bandwidth and isolate faults. Adversaries may delay, reorder, or drop messages, but cannot forge signatures under secure keys. Liveness holds once the network becomes timely for a sufficiently long window.

3.2 Admission Pipeline

The admission pipeline consists of decoding, stateless checks, witness checks, invariant checks, and tentative application. Stateless checks verify signature validity and nonce structure. Witness checks verify membership or non-membership against the current tentative accumulators. Invariant checks enforce application-level constraints such as non-negative balances, authorization policies, or rate limits. Upon success, the delta is applied to the tentative state and placed in the fiber mempool with a timestamp and cost weight.

3.3 Gossip Topology and Flow Control

Gossip is fiber-aware. Peers subscribe to fibers by interest and stake weight. Deltas are propagated with per-fiber rate limits and priority queues based on fee density and age. Backpressure is signaled via credit-based flow control: each peer advertises per-fiber capacity and grants credits for additional deltas only as it processes prior ones. This reduces congestion collapse during bursts and isolates spam to the targeted fiber.

3.4 Checkpoint BFT: Propose, Vote, Commit

Every T seconds the protocol advances a round. The leader samples the current roots from each fiber, bundles them into a vector R along with a DA summary, and broadcasts a proposal. Validators verify that R matches their locally derived roots or that discrepancies are bounded and explainable by admitted deltas plus DA samples. They then issue votes. A quorum certificate QC consisting of at least $2f+1$ signatures on the same R and round is sufficient to finalize the checkpoint. The scheme follows HotStuff-style linearity: propose, vote, and commit phases use threshold signatures to keep messages small.

3.5 Latency and Throughput Analysis

Checkpoint latency is approximately one network round-trip under synchrony plus local verification time, since payloads are small. Throughput scales with the number of active fibers and available DA bandwidth. The consensus layer processes constant-size vectors regardless of delta count, while admission and accumulator updates scale with per-fiber load. Choosing T balances confirmation latency and witness churn: smaller T lowers latency but increases checkpoint frequency and witness refresh rate.

3.6 Conflict Resolution Without Order

Conflicts such as double spends are resolved deterministically within a fiber using objective tiebreakers. A standard method binds each candidate to a verifiable random function output seeded by the signer's key and the conflict namespace. The admissible delta is the one with the smallest VRF output; others are excluded and provably non-members. This eliminates incentives to bribe a global sequencer for priority and collapses many forms of MEV.

3.7 Cross-Fiber Atomicity

Atomic composition across fibers uses two patterns. First, two-phase escrow: each fiber admits a conditional claim referencing the other fiber's claim and checkpoint number k . At checkpoint k both claims are promoted atomically if both are present and witnesses check out; otherwise both expire. Second, order-island mediation: a micro-sequencer produces a short transcript and a commitment that both fibers accept as a single delta. Both patterns avoid global ordering while preserving atomicity and auditability.

3.8 Pseudocode Sketch

Admission(fiber, delta): verify sig; check witnesses against roots; check invariants; apply tentative update; enqueue in fiber mempool. Propose(round): for each fiber collect root; assemble vector R and DA summary; broadcast proposal. Vote(round, R): verify fiber roots and DA; sign and send vote. Commit(QC): on $2f+1$ votes for R , finalize checkpoint; prune mempools; rotate leader.

4. Security and Data Availability

4.1 Adversary Assumptions

We consider Byzantine validators controlling up to f nodes, adaptive network scheduling, and rational clients. Cryptographic assumptions include hardness of discrete log or RSA as appropriate, unforgeability of signatures, and soundness of accumulator proofs. DA adversaries may try to withhold chunks or flood with garbage; sampling counters these with quantifiable failure probability.

4.2 Correctness Invariants

For each fiber, the following invariants hold across checkpoints: uniqueness of spent capabilities, monotonicity of effect accumulator, and consistency between tentative and finalized roots. Across fibers, escrow invariants ensure that conditional claims graduate or expire together.

4.3 Slashing Conditions and Evidence

Slashable offenses include: double-signing different vectors for the same round; endorsing a vector inconsistent with locally reconstructible roots and DA samples; and admitting deltas that violate invariants with publicly checkable evidence. Evidence bundles contain the signed messages and minimal witnesses needed for third parties to verify misbehavior.

4.4 Data Availability Sampling

Deltas are posted to an erasure-coded DA layer with total size N chunks and reconstruction threshold K . Each validator samples s random chunks per interval. If withholding exceeds a threshold, the probability that all samplers miss it decays exponentially in s . Checkpoints are considered valid only if a threshold of validators attach positive sample attestations. Parameters N , K , and s are chosen to bound undetected withholding below a target epsilon.

4.5 Censorship Resistance and Inclusion Deadlines

To mitigate sustained censorship by leaders, clients can submit deltas to multiple validators and to a public relay. Inclusion rules require that any valid delta that remains DA-available and uncontested be included within D checkpoints, or else a fault is attributable. This creates objective evidence of censorship for governance or slashing.

4.6 Denial-of-Service and Resource Accounting

Per-fiber fee markets price scarce resources such as accumulator updates and witness refreshes. Admission costs scale with proof sizes and update complexity. Validators enforce per-origin rate limits and use stateless filtering to drop obviously invalid deltas early. Order-island fibers charge higher fees reflecting their sequencing costs.

4.7 Key Management and Cryptographic Choices

RSA accumulators avoid trusted setup but have more expensive deletions; KZG commitments offer fast dynamic updates at the cost of a structured reference string. VRFs require secure key storage and rotation policies. The system should support pluggable cryptographic backends and versioned upgrades via checkpoints.

4.8 Liveness Proof Sketch

Under partial synchrony there exists a global stabilization time after which message delays are bounded. Given a correct leader and timely network, proposals carry valid root vectors and obtain $2f+1$ votes in one round-trip. Because deltas are admitted optimistically and DA is sampled continuously, any valid delta that is not in conflict will eventually be reflected in a finalized root.

5. Comparative Analysis

5.1 Blockchains

Blockchains enforce a single total order of transactions using longest-chain or BFT-based finality gadgets. Advantages include simple semantics and mature tooling. Costs include global serialization, fork choice complexity, and reorg risk. BAL differs by removing sequencing except for small checkpoint vectors, eliminating reorgs and reducing consensus bandwidth.

5.2 DAG Protocols

DAG systems encode partial order among transactions. They reduce leader bottlenecks but still reason about edges and topological order. Tip selection and conflict resolution introduce algorithmic complexity. BAL eschews graph semantics entirely for commutative fibers; verification reduces to checking accumulator witnesses against a signed vector.

5.3 Accumulator and Stateless Chains

Stateless or accumulator-based chains use cryptographic accumulators for compact proofs yet retain block sequences and ordering. BAL elevates accumulators to the ledger substrate, and consensus finalizes only the vector of roots. Thus accumulators move from optimization to core data model.

5.4 CRDT Ledgers

CRDT approaches show that many updates commute and can converge without centralized coordination, but typically assume cooperative settings or eventual consistency. BAL integrates CRDT-style deltas with BFT checkpoints and DA sampling to obtain strong safety with commutativity.

5.5 Rollups and Modular DA

BAL is complementary to modular designs. Fibers can post deltas to external DA networks and export succinct proofs to other ecosystems. Rollups can treat a fiber as their settlement scope, while BAL can serve as a high-throughput settlement layer for commutative workloads.

6. Use Cases

6.1 Payments

Model balances as capabilities identified by unique nonces. A transfer consumes a sender capability and creates a recipient capability with a fresh nonce. Witnesses prove membership of the sender's capability and non-membership of the new nonce in the spent-set. Conflicts are handled by VRF tie-breakers. Batched payrolls produce a single commitment per batch, lowering fees.

6.2 NFTs and Registries

Each collection maps to a fiber. Minting is a set insertion, transfer is a removal plus insertion, and burn is a removal with a tombstone. Atomic bundle trades use two-phase escrow where each involved fiber accepts conditional claims that graduate together at a checkpoint.

6.3 Social Identity and Messaging

Posts, reactions, and follows are naturally commutative. Revocations are signed removals. Clients subscribe to selected fibers and verify with small witnesses, enabling efficient feeds and portable identities without global traversal.

6.4 DeFi with Order-Islands

For order-dependent logic such as order books or liquidations, a local sequencer serializes events within an island and emits a single delta per interval for the parent fiber. Fair ordering or encrypted mempools can be applied locally. This confines MEV while preserving global commutativity.

6.5 Oracles and Bridges

Oracle committees sign price updates that are inserted as set elements. Consumers verify inclusion against the latest checkpoint. Bridges export accumulator witnesses to external chains, allowing verification without replicating BAL semantics.

6.6 Enterprise and Compliance

Enterprises can allocate private fibers with auditable checkpoints and selective disclosure. Auditors verify compliance statements via succinct proofs without exposure of unrelated data. This supports privacy-preserving attestations and regulated workflows.

7. Advantages

Scalability: workload partitions across fibers, removing global sequencing bottlenecks.

Consensus efficiency: checkpoint vectors are tiny, reducing bandwidth and leader load.

Finality: no chain to reorganize, so finalized membership is stable. Light clients: verify only the

fibers they touch with $O(1)$ proofs. MEV reduction: no ordering games outside order-islands.

8. Limitations

Expressiveness: applications must fit commutative primitives or use order-islands, which complicates some designs. Cryptographic assumptions: accumulators and VRFs introduce dependencies that must be maintainable across upgrades. Witness churn: clients must refresh proofs as accumulators evolve; this requires thoughtful wallet UX. Ecosystem maturity: new SDKs, explorers, and dev tooling are needed before production adoption.

9. Open Source Commitment

We will publish a reference node, light-client libraries, and sample applications under a permissive license such as Apache 2.0 or MIT. Governance will be open, with multiple independent implementations encouraged to avoid monoculture. Security reviews, formal verification of critical components, and transparent benchmarking will be prioritized.

10. Contact

For collaboration, research, and implementation feedback: X (Twitter) @xkal3b and Telegram @kb441.

11. Citation

Please cite: K. Barnhart, The Braid-Accumulator Ledger: A Novel Consensus and Data Structure Model for Decentralized Systems, 2025.

12. Conclusion

BAL reframes decentralized agreement around set membership instead of sequence. By combining scope-partitioned fibers, cryptographic accumulators, and periodic BFT checkpoints over vectors of roots, the model decouples throughput from global ordering and eliminates reorgs. Where strict order matters, BAL confines it to order-islands; elsewhere it allows commutative concurrency. The result is a third primitive that complements blockchains and DAGs while opening new design space for scalable, verifiable, and application-friendly decentralized systems.

13. Critical Perspectives and Open Questions

Witness churn economics: quantify wallet bandwidth for proof refresh across varying T and activity levels. Fee design: determine per-fiber fee markets that align incentives without fragmentation. Validator churn: specify reconfiguration policies that maintain safety during

membership changes. DA failures: analyze correlated withholding and recovery procedures. Cross-fiber atomicity: formalize failure modes and recovery for escrow and order-island mediation. Interoperability: standardize witness formats to ease bridging and light-client implementations.

14. References

- Nakamoto, S. 2008. Bitcoin: A Peer-to-Peer Electronic Cash System.
- Buterin, V. 2014. Ethereum: A Next-Generation Smart Contract Platform.
- Merkle, R. 1979. A Digital Signature Based on a Conventional Encryption Function.
- Castro, M., Liskov, B. 1999. Practical Byzantine Fault Tolerance.
- Yin, M., Malkhi, D., et al. 2019. HotStuff: BFT Consensus in the Lens of Blockchain.
- Micali, S., Rabin, M., Vadhan, S. 1999. Verifiable Random Functions.
- Kate, A., Zaverucha, G., Goldberg, I. 2010. Constant-Size Commitments to Polynomials and Their Applications.
- Benaloh, J., de Mare, M. 1993. One-Way Accumulators: A Decentralized Alternative to Digital Signatures.
- Nguyen, L. 2005. Accumulators from Bilinear Pairings and Applications.
- Bunz, B., et al. 2019. Proofs of Sequential Work and Stateless Cryptocurrency Discussions.
- Rocket Team. 2020. Avalanche Consensus Protocol.
- Popov, S. 2018. The Tangle.
- Baird, L. 2016. The Swirlds Hashgraph Consensus Algorithm.
- Shapiro, M., et al. 2011. A Comprehensive Study of CRDTs.
- Boneh, D., Bunz, B., Fisch, B. 2019. Batching Techniques for Accumulators.
- Gilad, Y., et al. 2017. Algorand: Scaling Byzantine Agreements.
- Buchman, E. 2016. Tendermint: Byzantine Fault Tolerance in the Age of Blockchains.
- Buterin, V., et al. 2020. Ethereum 2.0 and Gasper Discussions.
- Park, S., et al. 2021. Data Availability Sampling for Blockchains.
- Celestia Labs. 2021. Data Availability Whitepaper.